Serving Latency Notes

Abstract: request batching cuts tail latency in half.
Measured p95 drops from 820 ms to 410 ms.

中文摘要：批处理显著降低尾延迟。

https://example.com/latency-notes

Setup

We replay one week of recorded traffic across three batch sizes.